

Implications of Autonomy for the Expressiveness of Policy Routing

Nick Feamster, Ramesh Johari, *Member, IEEE*, and Hari Balakrishnan

Abstract—Thousands of competing autonomous systems must cooperate with each other to provide global Internet connectivity. Each autonomous system (AS) encodes various economic, business, and performance decisions in its routing policy. The current interdomain routing system enables each AS to express policy using *rankings* that determine how each router in the AS chooses among different routes to a destination, and *filters* that determine which routes are hidden from each neighboring AS. Because the Internet is composed of many independent, competing networks, the interdomain routing system should provide *autonomy*, allowing network operators to set their rankings independently, and to have no constraints on allowed filters. This paper studies routing protocol stability under these conditions. We first demonstrate that “next-hop rankings,” commonly used in practice, may not ensure routing stability. We then prove that, when providers can set rankings and filters autonomously, guaranteeing that the routing system will converge to a stable path assignment imposes strong restrictions on the rankings ASes are allowed to choose. We discuss the implications of these results for the future of interdomain routing.

Index Terms—Autonomy, BGP, interdomain routing, policy, safety.

I. INTRODUCTION

IN Internet routing, independently operated autonomous systems (ASes) must cooperate to exchange global information; nevertheless, this cooperation occurs in a landscape where these independent networks compete to provide Internet service. BGP facilitates this “competitive cooperation” by enabling network operators to express routing policies that are consistent with desired economic, business, and performance goals.

Ranking and *filtering* are the two main mechanisms that operators use to implement their policies. Ranking determines which of many possible routes to a destination should be used, thus providing an AS the flexibility to specify preferences over multiple candidate paths to a destination (e.g., specifying a primary




Fig. 1. Instability can arise when each AS independently specifies rankings [14], [22]. Each circle represents an AS. AS 0 is the destination. The listing of paths beside each node denotes a ranking over paths.

and a backup path). Filtering allows an AS to selectively advertise routes to some ASes and hide routes from others, thereby controlling which neighboring ASes send traffic over its infrastructure.

There are two important characteristics of policy routing: *autonomy* and *expressiveness*. Autonomy is the ability of each AS to set its rankings and filters independent of the others. Expressiveness refers to the flexibility that the routing protocol provides an operator for specifying rankings and filters. Ranking expressiveness determines what classes of rankings over routes are permitted by the protocol, while filtering expressiveness determines the range of route filters that are allowed.

The combination of expressiveness and autonomy has, in large part, been the reason for the success of BGP over the past decade. We contend that both autonomy and filtering expressiveness will be *requirements* for policy routing for the foreseeable future. Previous studies of routing stability assume that ASes are willing to compromise some degree of autonomy, filtering expressiveness, or both (see Section II). Autonomy, however, preserves each AS’s ability to set its policies without coordinating with any other AS. Filtering expressiveness gives an AS flexibility in how it establishes contracts with another AS, a task that should be unconstrained.

Ideally, an interdomain routing system should preserve autonomy, filtering expressiveness, and ranking expressiveness. However, the ability to specify highly expressive rankings comes at considerable cost to system robustness: as has been observed by Varadhan *et al.* and Griffin *et al.*, among others, if there are no constraints on the rankings that an AS can specify, BGP may violate safety (i.e., oscillate forever) [14], [22].

Example 1: Consider Fig. 1 [14], [22]. ASes 1, 2, and 3 each prefer the indirect path through their neighboring AS in the clockwise direction over the direct path to the destination, 0. All other paths are filtered. This configuration has no stable path assignment (i.e., a path assignment from which no node

Manuscript received October 25, 2005; revised May 14, 2006; approved by IEEE/ACM TRANSACTIONS ON NETWORKING Editor O. Bonaventure. This work was supported by the National Science Foundation (NSF) under Grants CNS-0225660 and CNS-0520241 and by a Cisco URP grant. The work of N. Feamster was supported in part by an NSF Graduate Research Fellowship. The work of R. Johari was supported in part by the Okawa Foundation. A preliminary version of this paper appeared in the Proceedings of the ACM SIGCOMM, August 2005.

N. Feamster is with the School of Computer Science, Georgia Institute of Technology, Atlanta, GA 30332 USA (e-mail: feamster@cc.gatech.edu).

R. Johari is with the Department of Management Science and Engineering, Stanford University, Stanford, CA 94305-4026 USA (e-mail: ramesh.johari@stanford.edu).

H. Balakrishnan is with the Electrical Engineering and Computer Science Department and the Computer Science and Artificial Intelligence Laboratory (CSAIL), Massachusetts Institute of Technology (MIT), Cambridge, MA 02139 USA (e-mail: hari@csail.mit.edu).

Digital Object Identifier 10.1109/TNET.2007.896531

would deviate). For example, consider the path assignment (10, 210, 30); in this case, AS 1 has a better path available, 130, so it switches paths. This switch breaks the path 210, causing AS 2 to switch to its second choice, path 20. The resulting path assignment, (130, 20, 30), is a permutation of the original path assignment: this time, AS 3 has the path 320 available, so it switches. This oscillation continues forever. ■

As the previous example suggests, full autonomy and expressiveness can have undesirable consequences. Routing protocol update messages should reflect actual reachability changes in the network topology or policy. Unfortunately, in BGP, conflicting policies can cause oscillations that produce endless streams of routing updates that are unrelated to changes in topology or policy. This instability creates numerous performance problems, may cause network partitions, and complicates diagnosis and debugging of problems in the routing system. Worse yet, a network operator has no way to guarantee that any given configuration of rankings and filters will not adversely interact with the policies of other ASes. In light of these issues, developing rigorous conditions on policy expressiveness that guarantee routing stability, while preserving autonomy, is crucial.

This paper explores the following question: if each AS retains complete autonomy and complete filtering expressiveness, how expressive can rankings be while guaranteeing stable routing? This question is important because ranking autonomy and filtering expressiveness reflect the realities of how ASes specify policies today, and little is known (beyond the results surveyed in Section II) about the tradeoffs between autonomy and expressiveness as far as routing stability is concerned, especially under filtering. In particular, our work is the first to develop necessary conditions for stability under realistic assumptions about autonomy and expressiveness and the first to derive necessary conditions for stability under filtering in policy routing.

We acknowledge that our requirements for complete autonomy and filtering expressiveness are strong; indeed, given the negative results presented in this paper, it may be reasonable to ask whether they are too strong. Nevertheless, we believe that the assumptions are not unnecessarily strong, because complete filtering expressiveness essentially reflects the realities of today's business landscape, where a provider has complete flexibility to form contracts (i.e., exchange routing information) with any other party. Thus, we believe that this paper studies routing stability guarantees from a reasonable point in the design space.

This paper makes three main contributions. First, in Section IV-A, we show that rankings based solely on the immediate next-hop AS en route to the destination may never reach a stable path assignment from an arbitrary initial state; i.e., next-hop rankings, which are common in practice, are *not safe*. Moreover, under unrestricted filtering, a routing system with next-hop rankings may have no stable path assignment. In addition to their operational implications, these results are also somewhat surprising, because next-hop rankings with no route filtering always have a stable path assignment. We also observe that, although rankings based on a globally consistent weighting

of paths are safe under filtering, even minor generalizations of the weighting function compromise safety (Section IV-B).

Second, we define a *dispute ring*, a special case of the “dispute wheel” (a group of nodes whose rankings have a particular form) of Griffin *et al.* [14], and show that any routing protocol that has a dispute ring is not safe under filtering (Section V). Using the dispute wheel concept, Griffin *et al.* showed a sufficient condition for safety, proving that if a routing system is unsafe then it must have a dispute wheel. In contrast, our result is the first known necessary condition for safety under filtering.

Third, we show that, providing for complete autonomy and filtering expressiveness, if providers directly choose their rankings, then the set of allowable rankings that guarantee safety is tightly constrained. In Section VI, we prove that any routing system that permits paths of length $n + 2$ to be ranked over paths of length n can have a dispute ring, and is thus unsafe under filtering. We also prove that any routing system that permits paths of length $n + 1$ to be ranked over paths of length n can have a dispute wheel. In summary, our results indicate that stable policy routing with provider autonomy and expressive filtering requires tight constraints on rankings.

Recent work has observed that routing protocols whose rankings are derived from a “strictly monotonic” algebra are guaranteed to converge [13], [15], [20], [21]; informally, a strictly monotonic algebra is one where a path has a higher cost (i.e., is less preferred) than any of its subpaths. In cases of unrestricted filtering, these strictly monotonic algebras represent a generalization of shortest paths routing, which is consistent with our results in Section VI. We discuss such mechanisms further in Section VII.

Our findings may be interpreted in several ways. The optimist will note that checking a set of rankings to ensure safety is trivial, because all it requires is that BGP routers modify the decision process to consult a route's “local preference” attribute only after considering its AS path length. The pessimist, however, may conclude that guaranteeing safe routing while preserving autonomy may yield constraints on expressiveness that are too strong. In either case, the results proved in this paper about the fundamental tradeoff between expressiveness and autonomy may help guide the design of stable interdomain routing protocols in the future. Section VII explores this tradeoff in further detail.

II. BACKGROUND

Because Internet routing is policy-based, and each AS has the flexibility to define its own policies independently of other ASes, the policies of one AS may interact with those of another to cause the protocol to oscillate. Seminal work by Varadhan *et al.* observed that policy-based interdomain routing protocols could oscillate and defined the concept of *safety* [22], [23]. Varadhan *et al.* also conjectured that routing systems that allow rankings other than those based on next-hop rankings or shortest path routing may be unsafe [22], [23]. In fact, in Section IV, we show that even routing systems that only allow next-hop rankings are not safe.




Fig. 2. Pairs of ASes may have different business relationships in different geographic regions.

Griffin *et al.* asked how expressive an autonomous, robust routing system can be [13]; we address this question in this paper. Varadhan *et al.* observed that a routing system with an acyclic topology will have at least one stable path assignment if participants can only express next-hop preferences [22], [23]. We show that when BGP's protocol dynamics are taken into account, restricting each AS to only next-hop rankings does *not* guarantee that the routing system will be safe (even though the routing system always has at least one stable path assignment).

Various previous work has studied *global* conditions to guarantee the safety of routing systems; global conditions presume that the routing system does not preserve local choice of rankings (i.e., ranking autonomy). Griffin *et al.* showed that, if the rankings of the ASes in a routing system do not form a *dispute wheel* (a concept that describes the global relationship between the rankings of a set of ASes), then the routing system is safe [14]. Griffin *et al.* also examined *robustness*, the property that safety is guaranteed even if arbitrary nodes or edges are removed from the graph. We view robustness as a special case of filtering: removing an edge can be achieved if the ASes incident to that edge filter all routes through that edge; removing a node entails having all ASes filter all routes through that node.

Gao and Rexford derived sufficient constraints on rankings, filtering, and global network topology to guarantee routing stability; they also observe that these constraints reflect today's common practice [8], [9]. They showed that if every AS considers each of its neighbors as either a customer, a provider, or a peer, and obeys certain local constraints on rankings and filtering, and if the routing system satisfies certain topological constraints, then BGP is stable.¹ However, their model does not incorporate ranking autonomy, because their proposed topological constraints are global. The example below illustrates why the restrictions of Gao and Rexford may sometimes be too strict.

Example 2: Consider Fig. 2; there are hundreds of similar real-world examples [2]. AS 1 and AS 2, two ISPs, peer in North America, but AS 1 buys service from AS 2 in Europe (in practice, this arrangement may occur if AS 1 does not have a European backbone). Suppose that AS 2 peers with AS 3 in North America, and AS 1 peers with AS 3 in Europe. This arrangement is precisely the “peer-provider” cycle that is prohibited by some of the convergence conditions of Gao and Rexford (e.g., Guideline A, [9]). This scenario mandates that AS 2 prefer the route to all destinations via AS 1 over AS 3, even destinations in North America. ■

The sufficient conditions of Griffin *et al.* require global knowledge of rankings [14]. Those of Gao and Rexford require

¹Griffin *et al.* noted that analogous sufficient conditions apply to iBGP with route reflection [17], although we show in previous work [3] that these conditions are unnecessarily strong.

global knowledge of the AS-level topology. Furthermore, the work of Griffin *et al.* does not consider the effects of filtering, and the conditions of Gao and Rexford restrict filtering severely. Our goal is to derive constraints that must hold on the configuration of a single AS *without any global knowledge*.

Griffin *et al.* also showed how to modify a BGP-like path vector protocol to detect the existence of a dispute wheel but left unspecified how the ASes should resolve the dispute wheel [16]. Machiraju and Katz defined a new global invariant for determining safety when at most one AS deviates from the conditions of Gao and Rexford [19]. Govindan *et al.* proposed a routing architecture where ASes coordinate their policies [11], [12] using a standardized policy specification language [1]. Jagard and Ramachandran presented global conditions that guarantee safety of routing systems that allow ASes to express only next-hop preferences over routes, and designed centralized and distributed algorithms to check these global conditions [18]. Gouda and Schneider study classes of routing protocol metrics for which each node in a routing tree has its most preferred path, but do not address routing stability [10]. In contrast to these studies of global conditions for safety, we study the conditions under which a policy-based interdomain routing protocol can be safe if it preserves the autonomy of each AS.

More recent work has attempted to design policy-based protocols that are guaranteed to converge without imposing any global conditions. Sobrinho described global relationships between preferences in an algebraic framework [20], [21], and incorporated several previous results (including those of both Griffin *et al.* [14] and Gao and Rexford [9]). He examined requirements for convergence and asserted that any vectoring protocol that preserves “strict monotonicity” is guaranteed to converge. Recent work on “metarouting” exploits this algebraic framework to allow designers to specify new routing protocols that are guaranteed to converge by requiring algebras that preserve strict monotonicity [15]. Strictly monotonic routing protocols may be viewed as extensions of shortest path routing to a more general algebraic setting; in this sense the design principles of metarouting are consistent with the characterization results of our paper. For further details, see Section VII.

III. ROUTING MODEL AND DEFINITIONS

We now define our routing model. After introducing some basic terminology, we formally define two notions of good behavior for routing protocols: *stability* and *safety*. Finally, we extend each of these two definitions to handle filtering expressiveness.

A. Preliminaries

We consider a model consisting of N ASes (nodes),² labeled $1, \dots, N$. Each of these nodes wishes to establish a path (defined below) to a single destination, labeled 0.

Definition 1 (Path): A path from i to j is a sequence of nodes $P = ii_1i_2 \dots i_mj$ with no repetition; i.e., such that $i_u \neq i_v$ if $u \neq v$, and $i_u \neq i, j$ for all u .

Note that we have defined the path to be loop-free, as BGP typically discards paths that have loops in the AS-level path.

²In this paper, we use the terms “AS” and “node” interchangeably.

Definition 1 also precludes the practice of AS-path prepending. However, the negative results of this paper continue to hold even if AS-path prepending is allowed.

We denote the number of hops in a path P as $\text{length}(P)$. In addition, given an AS k , we will write $k \in P$ if node k appears in P . For clarity, given a path P from i to j , we will often denote P by iPj ; furthermore, if P is a path from i to j , and Q is a path from j to k , then we will denote the concatenation of P and Q by $iPjQk$.

We denote the set of all paths from i to 0 (i.e., all paths on the complete graph) using the nodes $1, \dots, N$ by $\mathcal{P}_i^N$. Given the set of nodes $\{1, \dots, N\}$, each AS i will choose a ranking $\prec_i$ over the set of all paths $\mathcal{P}_i^N$, defined as follows.

Definition 2 (Ranking): A ranking $\prec_i$ for node i is a total ordering over the set of all paths $\mathcal{P}_i^N$; thus, given any two paths $P, Q \in \mathcal{P}_i^N$, either $P \prec_i Q$ (i prefers Q to P) or $P \succ_i Q$ (i prefers P to Q).

An AS may always choose the *empty path*, ε , which is equivalent to total disconnection from the destination node 0. Thus, we have $\varepsilon \in \mathcal{P}_i^N$ for all i and N . Furthermore, we assume that every AS strictly prefers connectivity to disconnectivity, so that $P \succ_i \varepsilon$ for all $P \in \mathcal{P}_i^N$.

All paths may not be available to node i , due to both topological constraints and filtering by other nodes. We will use $\mathcal{F}_i \subseteq \mathcal{P}_i^N$ to denote the set of paths actually available for use by node i . The empty path is always available; i.e., $\varepsilon \in \mathcal{F}_i$.

A routing system is specified by the rankings of the individual nodes, together with the paths available to the individual nodes. Observe that we have decoupled the “routing policy” of each AS i into two components: the rankings $\prec_i$ of AS i over route advertisements received; and a determination of which paths are filtered from other ASes. The filtering decisions of all nodes, together with physical constraints on the network, yield the sets $\mathcal{F}_1, \dots, \mathcal{F}_N$. We thus have the following formal definition of a routing system.

Definition 3 (Routing System): A routing system is a tuple $(N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$, where node i has ranking $\prec_i$ over the set $\mathcal{P}_i^N$, and $\mathcal{F}_i$ is the set of paths available to node i .

A routing system specifies the input to any interdomain routing protocol we might consider. Given this input, the protocol should converge to a “routing tree”: that is, an assignment of a path to each AS, such that the routes taken together form a tree rooted at 0. To formalize this notion, we must define path assignments and consistent paths.

Definition 4 (Path Assignment): A path assignment for the routing system $(N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$ is a vector of paths $\mathbf{P} = (P_1, \dots, P_N)$ such that, for all i , $P_i \in \mathcal{F}_i$.

Thus, a path assignment is an assignment of a feasible path to each node i , where feasibility is determined by the set of paths $\mathcal{F}_i$. Even though each node has a path assigned, these paths may not be *consistent*: node i may be assigned a path $P_i = ij\hat{P}_j0$, where j is the first node traversed on P_i , and where $\hat{P}_j$ is a path from j to 0. However, the path $\hat{P}_j$ may not be the same as the path P_j assigned to j in the path assignment $\mathbf{P}$; in fact, $\hat{P}_j$ may not even be in the set of feasible paths $\mathcal{F}_j$. For example, a node or link along the path $\hat{P}_j$ may experience a failure, causing the routing protocol to withdraw the path; if j has heard such a withdrawal but i has not, then it is possible that $P_i = ij\hat{P}_j0$

until node i learns that $\hat{P}_j$ no longer exists. To formally capture such situations, we define consistent paths and consistent path assignments.

Definition 5 (Consistent Path): Given a path assignment $\mathbf{P}$, a path $\hat{P}_i$ for node i is consistent with $\mathbf{P}$ if one of the following holds:

- 1) $\hat{P}_i = \varepsilon$; or
- 2) $\hat{P}_i = i0$; or
- 3) $\hat{P}_i = ijP_j0$, for some $j \neq i$.

Definition 6 (Consistent Path Assignment): A consistent path assignment for the routing system $(N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$ is a path assignment vector $\mathbf{P} = (P_1, \dots, P_N)$ such that for all i , P_i is consistent with $\mathbf{P}$.

A routing protocol where packets are forwarded based solely on destination should ultimately assign paths that are consistent with each other.

B. Stability and Safety

Informally, a path assignment is *stable* if it is consistent, and no node has a more preferred consistent path available.

Definition 7 (Stable Path Assignment): Given a routing system $(N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$, and a consistent path assignment $\mathbf{P}$, we say that $\mathbf{P}$ is *stable* if for all nodes i , and all other paths $\hat{P}_i \in \mathcal{F}_i$ that are consistent with $\mathbf{P}$, $\hat{P}_i \prec_i P_i$.

Definition 8 (Stable Routing System): The routing system $(N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$ is *stable* if there exists at least one stable path assignment $\mathbf{P}$.

The stability of a routing system does not indicate whether a routing protocol will converge *regardless* of the initial path assignment. The *safety* property, which states that a protocol eventually converges, regardless of the initial path assignment and ordering of the routing messages, captures this notion.

In defining safety, we will consider a simplified abstraction of BGP. We model the process by which nodes receive route advertisements from other nodes and subsequently update their own route decisions. We will consider a protocol dynamic where at each time step only a single AS is *activated*; when activated, an AS immediately processes all pending incoming route advertisements, and then makes a route decision. Formally, this model will translate into a path assignment sequence where exactly one node (the “activated” node) changes its route at any given time step.

A routing system is *safe* if no oscillation occurs regardless of the order in which nodes are activated.

Definition 9 (Fair Activation Sequence): The sequence $i_1, i_2, \dots$ is a fair activation sequence if each node $i = 1, \dots, N$ appears infinitely often in the sequence.

This definition of fair activation sequence is similar to that presented by Gao and Rexford [9], except that in our definition we only activate one node at a time. For our paper, this distinction is minor: we can interpret the Gao and Rexford dynamics as a model where outstanding routing messages may be in flight when a particular node is activated.

We now define our simplified model of the routing protocol dynamics: that is, starting from an initial path assignment $\mathbf{P}_0$, and given a fair activation sequence of nodes $i_1, i_2, \dots$, what is the resulting observed sequence of path assignments $\mathbf{P}_1, \mathbf{P}_2, \dots$? To formalize the dynamics of our model, we

Routing protocol dynamics

At time $t - 1$, the path assignment is P_{t-1} ; i.e., each node i has currently selected path $P_{i,t-1}$ to the destination 0. At time t :

- 1) A given node i_t is activated.
- 2) Node i_t updates its path to be the *most preferred path* (according to $\prec_{i_t}$) consistent with P_{t-1} . That is,
 - a) $P_{i_t,t} \in \mathcal{F}_{i_t}$ is consistent with P_{t-1} , and
 - b) $P_{i_t,t} \succ_{i_t} \hat{P}_{i_t} \forall \hat{P}_{i_t} \in \mathcal{F}_{i_t}$ consistent with P_{t-1} .
- 3) All other nodes leave their paths unchanged.

Fig. 3. The routing protocol dynamics, given an activation sequence $i_1, i_2, \dots$. The process starts from an initial path assignment P_0 .

consider an abstraction of the BGP decision process described in Fig. 3. At each time t , a node i_t is activated, and chooses its most preferred available path consistent with the path assignment P_{t-1} . All other nodes' paths remain unchanged. It is clear that this decision process yields a sequence of path assignments $P_1, P_2, \dots$. Note that this model assumes a very specific timing of message delivery and processing at the nodes; we discuss this issue further in Section V-A, where we compare our model to the dynamics of the Simple Path Vector Protocol (SPVP) model [14].

After any given activation step t , the overall path assignment P_t may not be consistent. Inconsistencies reflect the fact that a node only updates its path assignment in response to the receipt of a route advertisement. If, at time t_0 , a node i is using a path that traverses some other node j that has since changed paths, then node i would obviously continue to use (and advertise) that *inconsistent* path until it receives a routing update that reflects that the path through j has disappeared or changed. When activated, say, at time $t > t_0$, node i would discover that the path it was using was inconsistent with P_t and would then select its highest-ranked path that was consistent with P_t . The activation of a node at some time t corresponds to that node receiving all available routing information in the system up to that time.

With the definition of our protocol dynamics in hand, we can define protocol *safety*. Given a routing system and an activation sequence, we say that the system has converged if, after some finite time, the path assignment remains invariant for all future time. A protocol is *safe* if it converges to a stable path assignment, regardless of the initial path assignment and fair activation sequence.

Definition 10 (Safety): A routing system $(N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$ is safe if, for any initial path assignment P_0 and fair activation sequence $i_1, i_2, \dots$, there exists a finite T such that $P_t = P_T$ for all $t \geq T$.

Because the activation sequences are fair in the preceding definition, if a routing system converges to P_t , then the resulting path assignment to which the system converges must be both consistent and stable. If not, at least one node would change its path assignment eventually.

C. Filtering

We are interested in the stability and safety of systems that result when nodes are allowed to filter routes from other nodes. We thus require conditions stronger than stability and safety,

known as *stability under filtering* and *safety under filtering*. Informally, a routing system is stable (respectively, safe) under filtering if, under any choices of filters made by the ASes, the resulting routing system is always stable (respectively, safe).

Definition 11 (Stable Under Filtering): The routing system $(N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$ is stable under filtering if, for all choices of available paths $\hat{\mathcal{F}}_i \subseteq \mathcal{F}_i$ for $i = 1, \dots, N$, the routing system $(N, \prec_1, \dots, \prec_N, \hat{\mathcal{F}}_1, \dots, \hat{\mathcal{F}}_N)$ is stable.

Definition 12 (Safe Under Filtering): The routing system $(N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$ is safe under filtering if, for all choices of available paths $\hat{\mathcal{F}}_i \subseteq \mathcal{F}_i$ for $i = 1, \dots, N$, the routing system $(N, \prec_1, \dots, \prec_N, \hat{\mathcal{F}}_1, \dots, \hat{\mathcal{F}}_N)$ is safe.

We interpret these definitions as follows: The set of available paths $\mathcal{F}_i$ denotes the set of paths that are physically possible for node i to use, given the current network topology. Once all nodes have chosen their route filters, $\hat{\mathcal{F}}_i$ gives the set of paths that can ever be used by node i in route advertisements. Because we allow arbitrary choice of filters, the resulting routing system should be stable and safe regardless of the choices of $\hat{\mathcal{F}}_1, \dots, \hat{\mathcal{F}}_N$ that are made.

IV. RANKING CLASSES AND SAFETY

In this section, we study two natural ranking classes under which ASes retain autonomy in setting rankings over paths. First, in Section IV-A, we study the rankings where each AS is allowed to rank paths solely based on the immediate next-hop AS, called “next-hop rankings”. We show that 1) there are routing systems where each node has only a next-hop ranking that are not safe, and 2) even though all routing systems where nodes have next-hop rankings are stable, there exist some routing systems of this form that are not stable under filtering.

In Section IV-B, we study the properties of routing systems where each node is allowed to choose a weight for all its outgoing links, and rankings are derived from a “total” weight associated to each path. The total weight of a path is defined as the weight of the first link on that path, plus a discounted sum of the weights of all remaining links on that path. We show that if the discount factor is anything other than 1 (which corresponds to shortest path routing), then there exist weight configurations that yield a routing system that is not safe.

A. Next-Hop Rankings

One natural set of rankings for a routing system is one where each AS can express rankings over paths solely based on the next-hop AS in the path. Such a class of rankings makes sense because an AS establishes bilateral contracts with its immediate neighbors and, as such, will most often wish to configure its rankings based on the immediate next-hop AS en route to the destination. For example, an AS will typically prefer sending traffic via routes through its neighboring customer ASes over other ASes, since those customer ASes are paying based on traffic volume. We formally define next-hop rankings as follows:

Definition 13 (Next-Hop Ranking): Given N , $\prec_i$ is a next-hop ranking if, for all nodes j, k with i, j, k distinct, we have

$$ijP_{j0} \prec_i ikP_{k0} \Rightarrow ijP'_{j0} \prec_i ikP'_{k0} \quad (1)$$




Fig. 4. Next-hop rankings are not safe in this routing system. AS 1 prefers all paths through AS 3 over the direct path to the destination 0 (with ties broken deterministically) and prefers the direct path over all paths through AS 2. Similarly, AS 3 prefers all paths via AS 2, and so forth. (a) Routing system. (b) Activation sequence.

for all $P_j, P'_j \in \mathcal{P}_j^N$, and $P_k, P'_k \in \mathcal{P}_k^N$. (Here we interpret $\mathcal{P}_0^N = \{\varepsilon\}$.) Thus, $\prec_i$ ranks paths based only on the first hop of each path.

Such a restriction on policy would still be sufficiently rich to achieve most traffic engineering goals, since most policies are based on the immediate next-hop AS [4]. Additionally, this class of rankings is expressive enough for most current policy goals, because most current routing policies are dictated according to the AS's business relationship with its immediate neighbor. In this section, we show that while systems with next-hop rankings are generally stable, there exist examples that are unsafe, as well as systems that are unstable under filtering.

In the following proposition, we consider routing systems with next-hop rankings, provided that no filtering is employed. The proof is straightforward, using a construction due to Feigenbaum *et al.* [7].

Proposition 1: Suppose $(N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$ is a routing system such that $\prec_i$ is a next-hop ranking for each i , and $\mathcal{F}_i = \mathcal{P}_i$ for all i . Then there exists a stable path assignment $\mathbf{P}$ for this routing system.

We now show that there may exist $\hat{\mathcal{F}}_1 \dots \hat{\mathcal{F}}_N$, where $\hat{\mathcal{F}}_i \subseteq \mathcal{F}_i$ for all i , such that even though the system $(N, \prec_1 \dots \prec_N, \mathcal{F}_1 \dots \mathcal{F}_N)$ is stable, the filtered system $(N, \prec_1 \dots \prec_N, \hat{\mathcal{F}}_1 \dots \hat{\mathcal{F}}_N)$ is unstable. That is, there exist routing systems with next-hop rankings for which a stable path assignment exists, but introducing filtering can yield a system where no stable path assignment exists.

Observation 1: A routing system where each node has only a next-hop ranking may not be safe (even without filtering).

Example 3: Consider Fig. 4. In this example, each AS ranks every one of its neighboring ASes. For example, AS 1 prefers all paths that traverse AS 3 as the immediate next hop over all paths that traverse AS 0 as the immediate next hop, regardless of the number of ASes each path traverses; similarly, AS 1 prefers paths that traverse AS 0 as the immediate next hop over paths that traverse AS 2. Each AS readvertises its best path to the destination to all of its neighbors (i.e., the system has no filtering). Now consider the activation sequence in Fig. 4(b); if infinitely repeated, this activation sequence would be fair, and the routing system would oscillate forever. Thus, the routing system is not safe.

This system is not *safe*, but it is *stable*: for example, the path assignment (10, 210, 3210) is stable. Nodes 2 and 3 are using paths through their most preferred nodes. Node 1's most preferred node, node 3, is using a path that already goes through




Fig. 5. This routing system is stable without filtering but unstable under filtering. The figure shows a routing system with next-hop rankings and filtering that is equivalent to the unstable routing system with the rankings over paths shown in Fig. 1.

node 1, so node 1 is also using its most preferred consistent path. As every node is using its most preferred consistent path, no node will change paths when activated, so the path assignment is stable. ■

A routing system where each node has a next-hop ranking may not be safe, but Feigenbaum *et al.* showed that there is always guaranteed to be at least one stable path assignment for such routing systems [7]. However, allowing nodes to filter paths from each other can create routing systems for which there is no stable path assignment.

Observation 2: There exist routing systems with next-hop rankings for which a stable path assignment exists, but introducing filtering can yield a system where no stable path assignment exists.

Example 4: Consider Fig. 5. As before, each AS ranks every one of its neighboring ASes. Additionally, each AS may also declare arbitrary filtering policies. In this example, each AS (other than the destination) does *not* readvertise any indirect path to the destination. For example, AS 1 does not advertise the path 130 to AS 2, and thus the path 2130 is not available to AS 2. Formally, we define $\mathcal{F}_1 = \{130, 10\}$, $\mathcal{F}_2 = \{210, 20\}$, and $\mathcal{F}_3 = \{320, 30\}$.

The resulting routing system is equivalent to the system in Fig. 1, once the filtered paths are removed from each node's ranking. Thus, the filtered routing system is unstable by the same reasoning as that from the example shown in Fig. 1: for any path assignment in this routing system, at least one AS will have a higher ranked consistent path (and, hence, has an incentive to deviate from the path assignment). ■

The example demonstrates the complex interaction between filtering and rankings—a class of rankings that guarantees stability without filtering can be unstable under certain filtering conditions.

B. Edge Weight-Based Rankings

There exists at least one routing system that preserves autonomy and yet ensures safety under filtering: if each provider is allowed to choose edge weights for its outgoing links, and each provider ranks paths based on the sum of edge weights, the resulting “shortest paths” routing system is guaranteed to be safe [14]. Since this result holds for any $\mathcal{F}_1, \dots, \mathcal{F}_N$, any routing system built in this way is guaranteed to be safe under filtering. In this section, we will formulate a generalized model of such *edge weight-based rankings*, with both next-hop rankings and shortest path routing as special cases. Such rankings do not allow providers to directly specify their ranking; rather, the rankings of each provider are *derived* from the strategic choices

made by all providers, namely, the choices of outgoing link weights that each provider sets. This notion of “derived” rankings is a potentially useful method for ensuring autonomy in interdomain routing protocols; indeed, the metarouting framework can be viewed as a class of protocols based on derived rankings [15].

In [6], we consider a parameterized model of edge weight-based rankings. Total path weight is computed as the sum of the weight of the first hop on the path, and a parameter α times the total weight of the remainder of the path (where $\alpha \in [0, 1]$). Nodes prefer paths with lower total path weight. When $\alpha = 0$, we recover next-hop rankings; when $\alpha = 1$, we recover shortest path rankings. We show that the *only* edge weight-based ranking class that can guarantee safety under filtering, regardless of the weights chosen by each provider, is the class of shortest-path rankings.

V. DISPUTE WHEELS AND DISPUTE RINGS

Our goal is to study the classes of rankings for which the routing system is guaranteed to be safe under filtering. Griffin *et al.* have shown that checking whether a particular routing system is stable is NP-hard [14]. To simplify our study of safety, we introduce a useful concept developed by Griffin *et al.* [14], known as a *dispute wheel*. Informally, a dispute wheel gives a listing of nodes, and two path choices per node, such that the higher ranked path for each node uses the lower ranked path of its neighboring node on the wheel. The key property of a dispute wheel is that if a routing system oscillates, then it is possible to construct a dispute wheel where each node in the wheel selects its more preferred path (via the node in the clockwise direction) over its less preferred path. Griffin *et al.* showed that if a routing system with no filtering does not have a dispute wheel, then it is safe [14].

The dispute wheel is a useful concept because it allows us to analyze dynamic properties such as safety by simply looking at the rankings of each node in the routing system. In this section, we formally define a dispute wheel and show the relationship between the routing model of Griffin *et al.* in [14], which simulates messages being passed between nodes, to the model we use in this paper, which uses fair activation sequences. This relationship allows us to study safety in terms of the routing model in this paper. We then introduce a special type of dispute wheel called a *dispute ring* and show that, if any routing system has a dispute ring, then it is not safe under filtering. Finally, we relate dispute wheels to dispute rings and show that, although the presence of a dispute ring guarantees that a routing system is not safe under filtering, it does not necessarily imply that a routing system is not safe without filtering.

A. Dispute Wheels and Safety

We start by formally defining dispute wheels.

Definition 14 (Dispute Wheel): *Given a routing system $(N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$, a dispute wheel is a collection of distinct nodes $i_1, \dots, i_m$, called pivots, together with two sets of paths $P_1, \dots, P_m$ and $Q_1, \dots, Q_m$, such that the following conditions all hold (where we define $i_{m+1} = i_1$ for notational convenience):*

- 1) $P_k \in \mathcal{F}_{i_k}$ for all $k = 1, \dots, m$;




Fig. 6. Illustration of a dispute wheel. Dotted lines show preferred (indirect) paths to the destination. The nodes $i_1, \dots, i_m$ are pivots.

- 2) Q_k is a path from i_k to i_{k+1} for all $k = 1, \dots, m$;
- 3) The path $\hat{P}_k = i_k Q_k i_{k+1} P_{k+1} 0$ is feasible, i.e., $\hat{P}_k \in \mathcal{F}_{i_k}$;
- 4) $\hat{P}_k \succ_{i_k} i P_k 0$.

Thus, each node i_k prefers the path $i_k Q_k i_{k+1} P_{k+1} 0$ to the path $i_k P_k 0$, as shown in Fig. 6.

We now show that safety in the Simple Path Vector Protocol (SPVP) defined by Griffin *et al.* [14] implies safety in our model, which allows us to use dispute wheels to analyze safety.

Proposition 2: *Given a routing system, a fair activation sequence, and an initial path assignment P_0 , let $P_1, P_2, \dots$ be the resulting sequence of path assignments according to the dynamics described in Fig. 3. Then there exists a sequence of messages in the Simple Path Vector Protocol (SPVP) such that the same sequence of path assignments is observed.*

Thus, in particular, if a routing system is safe under SPVP, then it is safe according to Definition 10.

Proof Sketch: The key difference between SPVP and the dynamics we have defined is that SPVP is *asynchronous* (i.e., at any time step, messages may be in flight), so different nodes may have different assumptions about the global path assignment at any time. SPVP is *nondeterministic* with respect to the timing of messages; the delay between a routing update at node j and the receipt of the new route advertisement from node j at node i can be arbitrary. We use this fact to construct, inductively, a sequence of messages such that at time t , the current set of paths available to node i_t in SPVP corresponds exactly to P_{t-1} . Furthermore, we time the delivery of routing updates to node i_t in SPVP so that any updates that occurred since the last time i_t was activated arrive exactly at the start of time step t . In SPVP, this will initiate a routing update at node i_t , which corresponds exactly to the activation of i_t in our model (see Fig. 3).³

Thus, the sequence of path assignments seen in this realization of SPVP matches the sequence of path assignments seen in our dynamics. We conclude that if SPVP is guaranteed to be safe for the given routing system (i.e., if eventually no further routing updates occur, regardless of the initial path assignment), then the routing system is safe according to Definition 10 as well. ■

³A subtlety arises here, because in SPVP *each* of the routing updates that arrive to i_t at time t would trigger a routing update. However, in our candidate SPVP dynamics, we can set routing delays so that when nodes are activated after time t , the additional routing updates from node i_t are processed *instantaneously before* the last routing update that occurred at node i_t at time t . This ensures that any additional routing updates do not affect the observed sequence of path assignments, since the final path assignment at time t will still be P_t , as required.




Fig. 7. A routing system that is safe for any choice of filters.

Corollary 1: *If a routing system $(N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$ has no dispute wheel, then it is safe under filtering (and hence safe).*

Proof: Choose subsets $\hat{\mathcal{F}}_i \subseteq \mathcal{F}_i$. Then, any dispute wheel for the routing system $\hat{S} = (N, \prec_1, \dots, \prec_N, \hat{\mathcal{F}}_1, \dots, \hat{\mathcal{F}}_N)$ is also a dispute wheel for the original routing system $S = (N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$. Thus, the result follows from Proposition 2 and the results of [14]. ■

If no dispute wheel exists, the routing system is safe under filtering, but, unfortunately, this condition is not a necessary condition for safety, and thus not much can be said about a system that does have a dispute wheel. Furthermore, there exist routing systems that have a dispute wheel but which are safe under filtering.

Observation 3: *The existence of a dispute wheel does not imply that the routing system is unsafe, nor that the routing system is not safe under filtering.*

Example 5: See Fig. 7. The first two most preferred paths in each node's ranking form a dispute wheel, but the system is safe: the system converges to $\mathbf{P} = (10, 20, 30)$. Furthermore, *no combination of filters can create an oscillation*. The two-hop paths are not part of the stable path assignment, so filtering those paths has no effect on the protocol dynamics. Filtering a three-hop path would simply result in a node selecting the direct path to the destination, and the node would never deviate from that path. If one direct path is filtered, then the other two nodes will take direct paths to the destination and the node whose direct path is filtered will take its two-hop path. If two direct paths are filtered, then $\mathbf{P}$ is simply a chain to the destination: the node that has the direct path takes it, and the other two nodes will take two- and three-hop paths. ■

B. Dispute Rings and Safety

In this section, we extend the dispute wheel notion to understand the relationship between ranking expressiveness and safety under filtering. We define a relationship between rankings called a *dispute ring*, a special case of a dispute wheel where each node appears at most once. The dispute ring is a useful concept because it allows us to prove a *necessary* condition for safety under filtering.

Definition 15 (Dispute Ring): *A dispute ring is a dispute wheel—a collection of nodes $i_1, \dots, i_m$ and paths $P_1, \dots, P_m, Q_1, \dots, Q_m$ satisfying Definition 14—such that $m \geq 3$, and no node in the routing system appears more than once in the wheel.*




Fig. 8. Relationships between safety and dispute rings and wheels. Previous work showed that a routing system with no dispute wheel is safe [14]. Section V presents all other relationships shown in this figure. Corollary 1 states that any routing system without a dispute wheel is safe under filtering (and hence safe). Observation 3 notes that the reverse implications do not necessarily hold. Proposition 3 shows that any routing system with a dispute ring is unsafe under filtering.

Proposition 3: *If a routing system has a dispute ring, then it is not safe under filtering.*

Proof Sketch: We sketch the proof here; full details are available in [6]. Assume that a routing system has a dispute ring, defined by $i_1, \dots, i_m$, and paths $Q_1, \dots, Q_m, P_1, \dots, P_m$. Then, construct filters such that $\mathcal{F}_i$ contains *only* the paths in that dispute ring.

There exists at least one consistent path assignment $\mathbf{P}_t$ such that some pivot node i_{k-1} uses its most preferred path, $i_{k-1}Q_{k-1}i_kP_k0$, every other pivot node i_j uses path i_jP_j0 , and every other non-pivot node i uses its only available path consistent with this assignment. Then, the following activation sequence will result in an oscillation:

- 1) *Activate node i_k .* Node i_k then switches to its more preferred path, $i_kQ_ki_{k+1}P_{k+1}0$.
- 2) *Activate nodes along Q_{k-1} in reverse order, from the node immediately preceding i_k , to i_{k-1} .* All nodes along Q_{k-1} switch to the empty path, ε .
- 3) *Activate node i_{k-1} .* The path $i_{k-1}Q_{k-1}i_kP_k0$ is now inconsistent, so i_{k-1} must switch to the path $i_{k-1}P_{k-1}0$.
- 4) *Return to Step 1 with k replaced by $k+1$, and iterate again.*

By the fourth step of the iteration above, the new path assignment is “isomorphic” to the initial configuration: now node i_k is using the path $i_kQ_ki_{k+1}P_{k+1}0$, and every other pivot node i_j is using path i_jP_j0 . Thus, as this iteration repeats, the dynamics will ultimately reach node i_k once again with the original path assignment. ■

We emphasize that, for simplicity, we reduced the set of filters, $\mathcal{F}_i$, to include only the set of paths that are involved in an oscillation. We note that there will typically be more permissive sets $\mathcal{F}_i$ that will also result in oscillation, because the dispute ring is present in the underlying set of rankings. Our intent is to highlight the most basic case of filtering that can cause an oscillation, given the existence of a dispute ring.

Despite the fact that systems that are safe under filtering are guaranteed not to have a dispute ring, testing for a dispute ring is not sufficient to guarantee that the routing system is safe; see Figs. 10 and 11 in [6].

Fig. 8 summarizes the results of this section and how they relate to results from previous work [14].

VI. AUTONOMY AND SAFETY

In this section, we determine necessary and sufficient constraints on the allowable classes of rankings, such that if each AS autonomously sets its ranking while filtering is unrestricted,

the protocol is guaranteed to be safe. We do so by characterizing whether a routing system where rankings are independently specified by each AS can induce either a dispute ring or a dispute wheel.

Any protocol's configurable parameters implicitly constrain the rankings ASes can express. For example, in BGP, the set of protocol parameters is rich enough to allow providers to express essentially any possible ranking over paths. In Section VI-A, we axiomatically formulate two properties that should be satisfied by any protocol that respects autonomy: *permutation invariance* and *scale invariance*. The first requires the rankings allowed by the protocol to be independent of node labeling, while the second requires the allowed rankings to scale gracefully as nodes are added to the system. We abstract protocols satisfying these two conditions using the notion of an *autonomous ranking constraint* (ARC) function; such a function takes the ranking of a single AS as input, and accepts it if that ranking is allowed by the protocol. Observe that any protocol that respects the ability of ASes to autonomously choose rankings can be represented by a corresponding ARC function. Although we present the ARC function as a proof technique, this function could be implemented in practice to restrict the rankings that operators specify in operational networks.

In Section VI-B, we give two examples of such functions: the shortest hop count ARC function (which only accepts rankings where shorter paths are preferred to longer paths), and the next-hop ARC function (which only accepts next hop rankings). We then determine the class of ARC functions such that, as long as each node independently chooses an acceptable ranking, the resulting *global* routing system will be safe under filtering. In Section VI-C, we show that the only ARC functions that are safe under filtering are nearly equivalent to the shortest hop count ARC function.

A. ARC Functions

In this section, we define an *autonomous ranking constraint* (ARC) function, which serves as an abstraction of the protocol's constraints on allowed rankings over routes. We start by defining a local ranking constraint (RC) function, which takes as input the ranking of a single AS i , $\prec_i^N$, and determines whether that ranking is allowable.

Definition 16 (Local RC Function): Given N nodes, a local ranking constraint (RC) function $\pi(\prec_i)$ takes as input the ranking of a single AS i over all paths in $\mathcal{P}_i^N$, and returns “accept” if $\prec_i$ is allowed by π , and returns “reject” otherwise. If $\pi(\prec_i) = \text{“accept”}$, we will say that $\prec_i$ is π -accepted. If we are given a routing system $(N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$ where each $\prec_i$ is π -accepted, we will say the routing system is π -accepted.

Because we are restricting attention to protocols that respect the ability of ASes to choose rankings autonomously, a first condition that must be satisfied is that constraints on rankings should be “local”: that is, an AS should not face constraints on allowable rankings due to the rankings chosen by other ASes. For this reason, *local* RC functions act only on the ranking of a single AS. More generally, protocols might place system-wide constraints on the vector of rankings chosen by all ASes;

such protocols should be represented by “global” RC functions. Another generalization is to *derive* the preferences of an AS based on locally chosen parameters, as in the edge weight-based rankings of Section IV-B; RC functions, however, presume that providers directly choose their rankings over routes.

We now define two natural conditions any local RC function that preserves autonomy should satisfy. First, the local RC function's conditions on rankings should provide consistent answers to different ASes, regardless of the *labeling* of the ASes. That is, for the local RC function to preserve uniformity, each AS should be subject to the same constraints on routing policies, and those constraints should not depend on the particular assignment of AS numbers to ASes. For example, suppose the routing system consists of three ASes, and AS 1 has an accepted ranking where it prefers 1230 over 120, and 120 over 10. Then we expect the same ranking should be accepted, even if the labels of nodes are *permuted*. For example, suppose we permute the node labels that $1 \rightarrow 2$, $2 \rightarrow 3$, and $3 \rightarrow 1$. Then node 2 should also have an accepted ranking where it prefers 2310 over 230, and 230 over 20 (because 2310, 230, and 20 are the new paths that result after applying the permutation to 1230, 120, and 10, respectively). If this property were not satisfied, then the set of accepted rankings determined by a local RC function would depend on the global assignment of AS numbers to nodes, not on the characteristics of the individual rankings themselves. We call this notion *permutation invariance*; to define it precisely, we must proceed through a sequence of definitions, starting with *path permutation*.

Definition 17 (Path Permutation): Given N nodes, let σ be a permutation of the nodes $1, \dots, N$. Then σ induces a path permutation on any path $P = i_1 i_2 \dots i_m j$ from i to j , yielding a new path $\sigma(P) = \sigma(i) \sigma(i_1) \sigma(i_2) \dots \sigma(i_m) \sigma(j)$ from $\sigma(i)$ to $\sigma(j)$. We always define $\sigma(0) = 0$.

Definition 18 (Ranking Permutation): Given N nodes, let σ be a permutation of the nodes $1, \dots, N$. Then σ induces a ranking permutation on a ranking $\prec_i$ for node i over the paths in $\mathcal{P}_i^N$, yielding a new ranking $\sigma(\prec_i)$ over the paths in $\mathcal{P}_{\sigma(i)}^N$, as follows: If $P_1, P_2 \in \mathcal{P}_i^N$, and $P_1 \prec_i P_2$, then $\sigma(P_1) \sigma(\prec_i) \sigma(P_2)$ (where $\sigma(P_i)$ is the path permutation of path P_i under σ).

Note that a permutation does not modify the routing system in any substantive way, except to *relabel* the nodes, and to relabel the paths and rankings in a way that is consistent with the relabeling of nodes.

Definition 19 (Permutation Invariance): A local RC function π is permutation invariant if, given N and a ranking $\prec_i$ for an AS i over all paths in $\mathcal{P}_i^N$, the relation $\prec_i$ is π -accepted if and only if $\sigma(\prec_i)$ is π -accepted, for any permutation σ of $1, \dots, N$.

Second, a local RC function should specify conditions for acceptance or rejection of rankings that “scale” appropriately with the number of nodes in the system; we call this property *scale invariance*. Suppose, for example, that a local RC function accepts a ranking $\prec_i$ over $\mathcal{P}_i^N$, when N nodes are in the system. Now suppose that we add nodes to the system, so the total number of nodes is $\hat{N} > N$. As additional nodes are added to the system, additional paths become available as well, and each node i must specify its rankings over the larger set $\mathcal{P}_i^{\hat{N}}$. Informally, scale invariance of the local RC function requires that i should be able to “extend” the ranking $\prec_i$ to an accepted ranking over $\mathcal{P}_i^{\hat{N}}$,

without having to modify its existing ranking over $\mathcal{P}_i^N$; otherwise, the properties of allowed rankings would depend on the number of nodes in the global system.

To formalize this concept, we first define a subranking.

Definition 20 (Subranking): Given N nodes, let $\prec_i$ be a ranking for AS i over all paths in $\mathcal{P}_i^N$. Given $\hat{N} > N$, let $\hat{\prec}_i$ be a ranking for AS i over all paths in $\mathcal{P}_i^{\hat{N}}$. Note that $\mathcal{P}_i^N \subset \mathcal{P}_i^{\hat{N}}$. We say that $\prec_i$ is a subranking of $\hat{\prec}_i$ if $P_1 \prec_i P_2$ implies $P_1 \hat{\prec}_i P_2$, for all $P_1, P_2 \in \mathcal{P}_i^N$.

We now define scale invariance.

Definition 21 (Scale Invariance): A local RC function π is scale invariant if the following condition holds: given any π -accepted ranking $\prec_i$ for AS i over $\mathcal{P}_i^N$, and given any $\hat{N} > N$, there exists at least one π -accepted ranking $\hat{\prec}_i$ over $\mathcal{P}_i^{\hat{N}}$ that has $\prec_i$ as a subranking.

Permutation invariance guarantees that relabeling nodes does not affect allowed rankings; scale invariance ensures that even as the set of nodes in the network increases, the rankings over previously existing paths should remain valid. Local RC functions that satisfy both permutation invariance and scale invariance correspond to protocols that respect the ability of ASes to autonomously choose rankings; we call such functions *autonomous ranking constraint functions*.

Definition 22 (ARC Function): A local RC function is an autonomous ranking constraint (ARC) function if it is both permutation invariant and scale invariant.

We want to derive the conditions under which protocols are guaranteed to be safe under filtering. Given that we use an ARC function as an abstraction of the constraints placed by a protocol on rankings, we would thus like to characterize ARC functions that can ensure safety under filtering of the entire routing system (a global property). For this reason, we extend the definition of “safety under filtering” to cover local RC functions.

Definition 23: Let π be a local RC function. We say that π is safe under filtering if all π -accepted routing systems are safe under filtering.

B. Examples of ARC Functions

We now present two simple examples of ARC functions: the shortest hop count ARC function, which is guaranteed to be safe, but is not expressive; and the next hop ARC function, which is expressive, but not safe.

Example 6: Our first example is the *shortest hop count RC function*, π^{shc} . Given the number of nodes N , the RC function π^{shc} accepts a ranking $\prec_i$ for node i if and only if the relation $\prec_i$ strictly prefers shorter paths (in terms of hop count) over longer ones. Formally, it accepts $\prec_i$ if, for any two paths $P_i, \hat{P}_i \in \mathcal{P}_i^N$ such that $\text{length}(P_i) < \text{length}(\hat{P}_i)$, $P_i \succ_i \hat{P}_i$. Ties may be broken arbitrarily.

It is not hard to verify that π^{shc} is an ARC function. To check permutation invariance, note that if $\prec_i$ is allowed for node i , then of course for any permutation σ , the ranking $\sigma(\prec_i)$ will also be allowed for node $\sigma(i)$, as $\sigma(\prec_i)$ will also prefer shorter paths to longer paths. Scale invariance is natural: given any shortest hop count ranking $\prec_i$ over $\mathcal{P}_i^N$, and given $\hat{N} > N$, there obviously exists at least one shortest hop count ranking over $\mathcal{P}_i^{\hat{N}}$ that has $\prec_i$ as a subranking. ■

The RC function π^{shc} forces all providers to use shortest AS path length, effectively precluding each AS from having any policy expressiveness in choosing rankings (other than when breaking ties). A more flexible set of rankings is allowed by the next hop RC function of the next example.

Example 7: The next hop RC function, π^{nh} , accepts a ranking $\prec_i$ for node i if and only if $\prec_i$ satisfies (1) in Section IV-A; that is, if $\prec_i$ is a next hop ranking.

The RC function π^{nh} is clearly permutation invariant: if $\prec_i$ is a next hop ranking for node i , then clearly $\sigma(\prec_i)$ is a next hop ranking for node $\sigma(i)$. Furthermore, note that any next hop ranking $\prec_i$ is determined entirely by the rankings of node i over each possible next hop, together with tiebreaking choices among routes with the same next hop. Thus, for $\hat{N} > N$, $\prec_i$ can be extended to a next hop ranking over $\mathcal{P}_i^{\hat{N}}$, by extending node i 's rankings over each possible next hop, and determining tiebreaking rules for any routes with next hop $N + 1, \dots, \hat{N}$. We conclude that π^{nh} is scale invariant as well, and thus it is an ARC function.

The RC function π^{nh} grants greater flexibility in choosing routing policies than under the shortest hop count RC function, π^{shc} , albeit at some cost. With π^{nh} , each AS i will choose a next hop ranking $\prec_i$ without any *global* constraints on the composite vector of next hop rankings $(\prec_1, \dots, \prec_N)$ chosen by the nodes. We have shown earlier in Section IV-A that there exist configurations of next hop rankings that may not be stable or safe under filtering; thus, the ARC function π^{nh} can lead to a lack of safety. ■

Next, we use dispute rings and dispute wheels to characterize the class of ARC functions that are safe under filtering. We will prove that this class is closely related to the ARC function π^{shc} .

C. Impossibility Results

We prove two main results in this section. Informally, the first result can be stated as follows: suppose we are given an ARC function and an accepted ranking such that some n hop path is *less preferred* (i.e., ranked lower) than another path of length at least $n + 2$ hops. Then, we can construct an accepted routing system with a dispute ring; i.e., one that is not safe under filtering. The second result states that if some n -hop path is *less preferred* than another path of length at least $n + 1$ hops, then there exists a routing system with a dispute wheel (though not necessarily a dispute ring). Note that this result is weaker than our first result, because a dispute wheel does not necessarily imply that the system is not safe under filtering.

We interpret these results as follows: if we are searching for ARC functions that are safe under filtering, we are very nearly restricted to considering only the shortest hop count ARC function, because all paths of n hops *must be more preferred* than paths of at least $n + 2$ hops to guarantee safety under filtering, and all paths of n hops must be more preferred than paths of at least $n + 1$ hops to prevent dispute wheels.

Our first lemma, which is crucial to proving both of our results, uses permutation invariance to construct a dispute wheel from a single node's rankings. We use a permutation to “replicate” pieces of the dispute wheel until the entire wheel is completed.




Fig. 9. Dispute wheel construction for Lemma 1.

To state the lemma, we must define the *period* of a node with respect to a permutation, as well as the period of a permutation. Given a permutation σ on the nodes $1, \dots, N$, let σ^k denote the permutation that results when σ is applied k times; e.g., $\sigma^2(j) = \sigma(\sigma(j))$, where σ^0 is defined to be the identity permutation.

Definition 24 (Period): Given a permutation σ on the nodes $1, \dots, N$, we define the period of i under σ as $\text{period}_i(\sigma) = \min\{k \geq 1 : \sigma^k(i) = i\}$.

Thus, the period of i is the minimum number of applications of σ required to return to i .

Definition 25 (Permutation Period): Given a permutation σ on the nodes $1, \dots, N$, we define the period of the permutation σ as $\text{period}(\sigma) = \min\{k \geq 1 : \sigma^k(i) = i \text{ for all } i\}$.

Thus, $\text{period}(\sigma)$ is the minimum number of applications of σ required to recover the original node labeling, and can be computed as the least common multiple of $\text{period}_i(\sigma)$, for $1 \leq i \leq N$.

The following result establishes the conditions under which we can apply a permutation to a π -accepted ranking to obtain a dispute wheel. We use this lemma as a building block for both of the theorems in this section.

Lemma 1: Let π be an ARC function. Suppose there exists a node i with a ranking $\prec_i$ over $\mathcal{P}_i^N$, two paths $R_i, \hat{P}_i \in \mathcal{P}_i^N$, and a permutation σ on $1, \dots, N$ such that: 1) $\prec_i$ is π -accepted; 2) $R_i \succ_i \hat{P}_i$; 3) $\text{period}_i(\sigma) = \text{period}(\sigma)$; and 4) there exists a path $\hat{Q}_i$ from i to $\sigma(i)$ such that

$$R_i = i\hat{Q}_i\sigma(i)\sigma(\hat{P}_i)0. \quad (2)$$

Then there exists a π -accepted routing system with a dispute wheel.

This dispute wheel is defined by pivot nodes $i_1, \dots, i_m$, and paths $P_1, \dots, P_m$ and $Q_1, \dots, Q_m$, where $m = \text{period}(\sigma)$, and where for $k = 1, \dots, m$, we have $i_k = \sigma^{k-1}(i)$, $P_k = \sigma^{k-1}(\hat{P}_i)$, and $Q_k = \sigma^{k-1}(\hat{Q}_i)$.

Proof: Refer to Fig. 9. The key idea of the proof is that, since $\text{period}_i(\sigma) = \text{period}(\sigma)$, we can repeatedly apply σ to the paths $\hat{Q}_i$ and $\hat{P}_i$ and apply permutation invariance to construct a π -accepted routing system with a dispute wheel.

Let $m = \text{period}(\sigma)$. Define the sequence $i_1, i_2, \dots, i_m$ by $i_k = \sigma^{k-1}(i)$ for $k = 1, \dots, m$. Since $\text{period}(\sigma) = \text{period}_i(\sigma)$, the nodes $i_1, \dots, i_m$ are all distinct. For $k = 1, \dots, m$, define $\prec_{i_k} = \sigma^{k-1}(\prec_i)$; since the nodes

$i_1, \dots, i_m$ are all distinct, this assignment of rankings to nodes is well defined (i.e., no node is assigned two different rankings). By permutation invariance, since $\prec_i$ is π -accepted, we conclude $\prec_{i_k}$ is π -accepted for all k . For all other nodes j , choose any π -accepted ranking $\prec_j$. Let $\mathcal{F}_j = \mathcal{P}_j^N$ for all nodes j .

This permutation defines a π -accepted routing system $(N, \prec_1, \dots, \prec_N, \mathcal{F}_1, \dots, \mathcal{F}_N)$. We now construct a dispute wheel for this system. Define $Q_k = \sigma^{k-1}(\hat{Q}_i)$, and $P_k = \sigma^{k-1}(\hat{P}_i)$, for $k = 1, \dots, m$. We claim that these definitions yield a dispute wheel.

Since $\mathcal{F}_j = \mathcal{P}_j^N$ for all j , all paths are feasible. Next, since $\hat{Q}_i$ is a path from $i_1 = i$ to $i_2 = \sigma(i)$, we conclude that Q_k is a path from i_k to i_{k+1} for all k (where we define $i_{m+1} = i_1$). We now observe that

$$\begin{aligned} \sigma^{k-1}(R_i) &= \sigma^{k-1}(i)\sigma^{k-1}(\hat{Q}_i)\sigma^k(i)\sigma^k(\hat{P}_i)0 \\ &= i_k Q_k i_{k+1} P_{k+1} 0. \end{aligned}$$

Finally, since $\prec_{i_k} = \sigma^{k-1}(\prec_i)$ and $R_i \succ_i \hat{P}_i$, we have $\sigma^{k-1}(R_i) \succ_{i_k} \sigma^{k-1}(\hat{P}_i)$. Using the preceding derivation and the fact that $P_k = \sigma^{k-1}(\hat{P}_i)$, we conclude that $i_k Q_k i_{k+1} P_{k+1} 0 \succ_{i_k} i_k P_k 0$, as required.

Thus, we have established that $i_1, \dots, i_m$, together with $Q_1, \dots, Q_m$ and $P_1, \dots, P_m$, constitute a dispute wheel. ■

The preceding lemma reduces the search for a dispute wheel to a search for a permutation and a π -accepted ranking with the stated properties. Observe from (2) that the permutation σ maps the path $\hat{P}_i$ into the “tail” of the path R_i ; in applying Lemma 1, we will construct a partial permutation by mapping a path $\hat{P}_i$ into the “tail” of R_i as in (2), and then we will complete the permutation by adding nodes to the system if necessary so that $\text{period}_i(\sigma) = \text{period}(\sigma)$. We use this approach to prove two theorems; the first states that if an ARC function accepts at least one ranking that prefers an n -hop path less than a path of at least $n + 2$ hops, then the ARC function is not safe under filtering.

Theorem 1: Let π be an ARC function. Suppose there exist a node i with π -accepted ranking $\prec_i$ and two paths $R_i, \hat{P}_i \in \mathcal{P}_i^N$, such that $\text{length}(R_i) > \text{length}(\hat{P}_i) + 1$ and $R_i \succ_i \hat{P}_i$. Then, π is not safe under filtering.

Proof: The proof relies on Lemma 1 to build a dispute wheel. First, using scale invariance of the ARC function, we show that the stated conditions of the theorem ensure that there exist two paths $R'_i, \hat{P}'_i$ such that: $\text{length}(R'_i) \geq \text{length}(\hat{P}'_i) + 1$; R'_i is more preferred than $\hat{P}'_i$ for some π -accepted ranking; and R'_i and $\hat{P}'_i$ have no nodes in common, other than i and 0. Lemma 2 then completes the proof of the theorem through two steps: first, once we have found the paths R'_i and $\hat{P}'_i$, we use them to build a permutation σ such that the conditions of Lemma 1 are satisfied; and second, we show that the dispute wheel given by Lemma 1 is in fact a dispute ring, by checking that no nodes are repeated around the wheel.

We first construct the paths R'_i and $\hat{P}'_i$ as described in the previous paragraph. Let $i, \prec_i, R_i$, and $\hat{P}_i$ be given as in the theorem. Let $\ell = \text{length}(\hat{P}_i)$; i.e., $\hat{P}_i = iu_1u_2 \dots u_{\ell-1}0$. We add ℓ new nodes to the routing system, and label them $v_1, \dots, v_\ell$; let $N' = N + \ell$. By scale invariance, there exists a π -accepted ranking $\prec_{i'}^{N'}$ on the set of paths $\mathcal{P}_i^{N'}$ with $\prec_i$ as a subranking. For such a ranking $\prec_{i'}^{N'}$ we have $R_i \succ_{i'}^{N'} \hat{P}_i$.

But now consider the path $T_i = iv_1 \dots v_\ell 0$; note that $\text{length}(T_i) = \ell + 1$. Since $R_i \succ_i^{N'} \hat{P}_i$, either $T_i \succ_i^{N'} \hat{P}_i$, or $R_i \succ_i^{N'} T_i$. In the former case, let $R'_i = T_i$, $\hat{P}'_i = \hat{P}_i$; and in the latter case, let $R'_i = R_i$, and $\hat{P}'_i = T_i$. Then $\text{length}(R'_i) \geq \text{length}(\hat{P}'_i) + 1$, $R'_i \succ_i^{N'} \hat{P}'_i$, and R'_i and $\hat{P}'_i$ have no nodes in common other than i and 0 .

The following lemma uses Lemma 1 to construct a dispute wheel.

Lemma 2: *Let π be an ARC function. Suppose there exists a node i with π -accepted ranking $\prec_i$ over $\mathcal{P}_i^N$, and two paths $R_i, \hat{P}_i \in \mathcal{P}_i^N$ such that:*

- 1) $\text{length}(R_i) \geq \text{length}(\hat{P}_i) + 1$;
- 2) $R_i \succ_i \hat{P}_i$; and
- 3) R_i and $\hat{P}_i$ have no nodes in common other than i and 0 .

Then there exists a π -accepted routing system for which there exists a dispute ring.

Proof of Lemma: The proof of this lemma proceeds by using scale invariance: we add enough new nodes to the system to allow us to build a permutation such that the conditions of Lemma 1 are satisfied. The key insight is that we initially construct the permutation σ by mapping the path $\hat{P}_i$ into the “tail” of the path R_i . We then add enough nodes so that when we complete the definition of σ , we have $\text{period}_i(\sigma) = \text{period}(\sigma)$.

Let $\text{length}(\hat{P}_i) = n$, and let $h = \text{length}(R_i) - n$; note that, by Condition 1 in the statement of the lemma, we know $h \geq 1$. Define $i_1 = i$. We label the nodes so that $\hat{P}_i = i_1 i_2 \dots i_n 0$, and $R_i = i_1 x_1 x_2 \dots x_{h-1} \hat{i}_1 \hat{i}_2 \dots \hat{i}_n 0$. We want to define a permutation σ that will map the path $i_1 \dots i_n 0$ to the tail of R_i , i.e., to the path $\hat{i}_1 \dots \hat{i}_n 0$. However, this does not completely define a permutation, so we must add additional nodes to ensure that $\text{period}_i(\sigma) = \text{period}(\sigma)$.

We add $2(h-1) + n$ additional nodes to the system, labeled $\hat{x}_1, \dots, \hat{x}_{h-1}$, and $i'_1, \dots, i'_n, x'_1, \dots, x'_{h-1}$. By scale invariance, we know there exists at least one π -accepted ranking $\prec_i$ over all paths using this larger set of nodes, such that $\prec_i$ has $\prec_i$ as a subranking. In particular, since $R_i \succ_i \hat{P}_i$, we have $R_i \succ_i \hat{P}_i$. We now define a permutation σ according to the following maps:

$$\begin{aligned} i_k &\rightarrow \hat{i}_k \rightarrow i'_k \rightarrow i_k, & k = 1, \dots, n; \\ x_k &\rightarrow \hat{x}_k \rightarrow x'_k \rightarrow x_k, & k = 1, \dots, h-1. \end{aligned}$$

That is, $\sigma(i_k) = \hat{i}_k$, $\sigma(\hat{i}_k) = i'_k$, etc. For all nodes j not listed, we define $\sigma(j) = j$. Note that the period of σ is $\text{period}(\sigma) = 3$, and of course $\text{period}_i(\sigma) = \text{period}_{i_1}(\sigma) = 3 = \text{period}(\sigma)$. Finally, note that by definition of σ , we have $R_i = i \hat{Q}_i \sigma(i) \sigma(\hat{P}_i) 0$, where $\hat{Q}_i = i x_1 \dots x_{h-1} \hat{i}_1$.

Thus, the conditions of Lemma 1 have been satisfied by the ranking $\prec_i$, the paths R_i and $\hat{P}_i$, and the permutation σ ; so we know there exists a π -accepted routing system for which there exists a dispute wheel. To complete the proof, we need only check that the dispute wheel is a dispute ring. Note that the wheel has three pivot nodes. Furthermore, to check that no nodes are repeated around the wheel, we simply enumerate the elements of our dispute wheel: $\hat{Q}_i = i x_1 \dots x_{h-1} \hat{i}_1$; $\sigma(\hat{Q}_i) = \hat{i}_1 \hat{x}_1 \dots \hat{x}_{h-1} i'_1$; $\sigma^2(\hat{Q}_i) = i'_1 x'_1 \dots x'_{h-1} \hat{i}_1$; $\hat{P}_i = i_1 \dots i_n 0$; $\sigma(\hat{P}_i) = \hat{i}_1 \dots \hat{i}_n 0$; and $\sigma^2(\hat{P}_i) = i'_1 \dots i'_n 0$. It is straightforward to check that these paths constitute a dispute ring: in




Fig. 10. Dispute ring construction for Lemma 2.

Fig. 10, note that the dispute wheel constructed from these paths has no repeated nodes. ■

Lemma 2 completes the proof of the theorem: we have shown that if some π -accepted ranking exists satisfying the conditions of the theorem, then using only permutation invariance and scale invariance we can build a π -accepted routing system with a dispute ring. This routing system is then unsafe under filtering, by Proposition 3. ■

The preceding theorem suggests that ARC functions that are safe under filtering are closely related to the shortest hop count ARC function, because no rankings can be accepted where n hop paths are less preferred than $(n+k)$ -hop paths, for $k \geq 2$. The next theorem draws this relationship even closer, by proving that there exists a dispute wheel if an ARC function accepts any ranking where an n -hop path is less preferred than an $(n+1)$ -hop path.

Theorem 2: *Let π be an ARC function. Suppose there exist a node i with π -accepted ranking $\prec_i$ and two paths $R_i, \hat{P}_i \in \mathcal{P}_i^N$, such that $\text{length}(R_i) = \text{length}(\hat{P}_i) + 1$ and $R_i \succ_i \hat{P}_i$. Then there exists a π -accepted routing system with a dispute wheel.*

Proof Sketch: Details are in the technical report [5]. Our approach is to map the path $\hat{P}_i$ into the “tail” of the path R_i , which partially defines a permutation σ . We then show how to add nodes to the system and complete the permutation σ so that $\text{period}_i(\sigma) = \text{period}(\sigma)$. We finally apply Lemma 1 to conclude there exists a π -accepted routing system with a dispute wheel. ■

The preceding results should not be interpreted as suggesting that we cannot find a routing system that is safe under filtering, where nodes prefer $(n+k)$ -hop paths over n -hop paths. Indeed, from Fig. 7, there are routing systems where nodes prefer three-hop paths over one-hop paths, and yet the system is safe under filtering. However, checking whether such systems are safe under filtering requires global verification; the theorems we have presented suggest that safety under filtering cannot be guaranteed through local verification alone, if some nodes are allowed to prefer $(n+k)$ -hop paths over n -hop paths.

Furthermore, the two results in this section highlight the importance of dispute rings. Theorem 1 gives the strong result that an ARC function that allows some $(n+k)$ -hop path to be more preferred than an n -hop path cannot guarantee safety under filtering, if $k \geq 2$. Theorem 2 only guarantees the existence of a

dispute wheel if an ARC function that allows some $(n+1)$ -hop path to be more preferred than an n -hop path—we cannot draw conclusions regarding the stability or safety of a routing system on the basis of the existence of a dispute wheel (see Fig. 7).

VII. SUMMARY

This paper explored the fundamental tradeoff between the expressiveness of rankings and routing safety, presuming that each AS retains complete autonomy and filtering expressiveness, and presented the first study of the effects of filtering on safety. We make the following contributions: 1) We showed that next-hop rankings are not safe; we also observed that although rankings based on a globally consistent weighting of paths are safe under filtering, even minor generalizations of the weighting function compromise safety. 2) We defined a *dispute ring* and show that any routing system that has a dispute ring is not safe under filtering. Our results are the first necessary conditions concerning safety. 3) We showed that, providing for complete autonomy and filtering expressiveness, if providers directly choose their rankings, the class of allowable rankings that guarantee safety is strongly constrained.

Although the results of Section VI appear to tightly constrain allowable rankings, Section IV-B suggests one possible direction forward: there we note that routing using preferences derived from edge weights is guaranteed to be stable. Suppose each AS ranks paths based on the sum of edge weights to the destination and adjusts weights on its incident outgoing edges to neighboring ASes. Rankings would then be derived from the total path cost, but an AS might still retain enough flexibility to control the next-hop AS en route to the destination. More generally, rankings may be built based on consistent edge labels aggregated using a strictly monotonic algebra [15], [20], [21]; this class of routing protocols using “derived” rankings is guaranteed to converge. Such an approach could ensure that the protocol is safe on short timescales, while allowing “policy disputes” to occur on longer timescales, out-of-band from the routing protocol. Of course, we must still explore whether this apparently more restrictive language could still implement the policies that operators want to express. Furthermore, a more restrictive policy language would guarantee safety, but would likely cause routing to oscillate on a slower timescale as operators observe the routing protocol converging to undesirable paths. It is our position that this design decision is exactly the right one: the routing protocol *should* converge on a fast timescale and accurately reflect network topology, while policy conflicts should be resolved on slower, “human” timescales.

REFERENCES

- [1] C. Alaettinoglu *et al.*, Routing Policy Specification Language (RPSL), RFC 2622, Jun. 1999.
- [2] Randy Bush, May 2004, private communication.
- [3] N. Feamster and H. Balakrishnan, “Correctness properties for Internet routing,” presented at the 43rd Annu. Allerton Conf. Communication, Control, and Computing, Monticello, IL, Sep. 2005.

- [4] N. Feamster, J. Borkenhagen, and J. Rexford, “Guidelines for interdomain traffic engineering,” *ACM Comput. Commun. Rev.*, vol. 33, no. 5, pp. 19–30, Oct. 2003.
- [5] N. Feamster, R. Johari, and H. Balakrishnan, “Stable policy routing with provider independence,” Mass. Inst. of Technol., Tech. Rep. MIT-LCS-TR-981, Feb. 2005 [Online]. Available: <http://www.lcs.mit.edu/publications/specpub.php?id=1715>
- [6] N. Feamster, R. Johari, and H. Balakrishnan, “The implications of autonomy for stable policy routing,” in *Proc. ACM SIGCOMM*, Philadelphia, PA, Aug. 2005, pp. 25–36.
- [7] J. Feigenbaum, R. Sami, and S. Shenker, “Mechanism design for policy routing,” in *Proc. ACM Symp. Principles of Distributed Computing*, 2004, pp. 11–20.
- [8] L. Gao, T. G. Griffin, and J. Rexford, “Inherently safe backup routing with BGP,” in *Proc. IEEE INFOCOM*, Anchorage, AK, Apr. 2001, pp. 547–556.
- [9] L. Gao and J. Rexford, “Stable Internet routing without global coordination,” *IEEE/ACM Trans. Networking*, vol. 9, no. 6, pp. 681–692, Dec. 2001.
- [10] M. Gouda and M. Schneider, “Maximizable routing metrics,” *IEEE/ACM Trans. Networking*, vol. 11, no. 4, pp. 663–675, Aug. 2003.
- [11] R. Govindan, C. Alaettinoglu, G. Eddy, D. Kessens, S. Kumar, and W. Lee, “An architecture for stable, analyzable Internet routing,” *IEEE Network*, vol. 13, no. 1, pp. 29–35, Jan./Feb. 1999.
- [12] R. Govindan, C. Alaettinoglu, K. Varadhan, and D. Estrin, “Route servers for inter-domain routing,” *Computer Networks and ISDN Systems*, vol. 30, pp. 1157–1174, 1998.
- [13] T. Griffin, A. Jaggard, and V. Ramachandran, “Design principles of policy languages for path vector protocols,” in *Proc. ACM SIGCOMM*, Karlsruhe, Germany, Aug. 2003, pp. 61–72.
- [14] T. Griffin, F. B. Shepherd, and G. Wilfong, “The stable paths problem and interdomain routing,” *IEEE/ACM Trans. Networking*, vol. 10, no. 1, pp. 232–243, Feb. 2002.
- [15] T. Griffin and J. L. Sobrinho, “Metarouting,” in *Proc. ACM SIGCOMM*, Philadelphia, PA, Aug. 2005, pp. 1–12.
- [16] T. Griffin and G. Wilfong, “A safe path vector protocol,” in *Proc. IEEE INFOCOM*, Tel Aviv, Israel, Mar. 2000, pp. 490–499.
- [17] T. Griffin and G. Wilfong, “On the correctness of IBGP configuration,” in *Proc. ACM SIGCOMM*, Pittsburgh, PA, Aug. 2002, pp. 17–29.
- [18] A. D. Jaggard and V. Ramachandran, “Robustness of class-based path vector systems,” in *IEEE Int. Conf. Network Protocols (ICNP)*, Berlin, Germany, Oct. 2004, pp. 84–93.
- [19] S. Machiraju and R. Katz, “Verifying global invariants in multi-provider distributed systems,” in *Proc. 3rd ACM Workshop on Hot Topics in Networks (Hotnets-III)*, San Diego, CA, Nov. 2004, pp. 149–154.
- [20] J. L. Sobrinho, “Network routing with path vector protocols: Theory and applications,” in *Proc. ACM SIGCOMM*, Karlsruhe, Germany, Aug. 2003, pp. 49–60.
- [21] J. L. Sobrinho, “An algebraic theory of dynamic network routing,” *IEEE/ACM Trans. Networking*, vol. 13, no. 5, pp. 1160–1173, Oct. 2005.
- [22] K. Varadhan, R. Govindan, and D. Estrin, “Persistent route oscillations in inter-domain routing,” USC/ISI, Tech. Rep. 96-631, Feb. 1996.
- [23] K. Varadhan, R. Govindan, and D. Estrin, “Persistent route oscillations in inter-domain routing,” *Computer Networks*, vol. 32, no. 1, pp. 1–16, 2000.



Nick Feamster received the S.B. and M.Eng. degrees in electrical engineering and computer science in 2000 and 2001, respectively, and the Ph.D. degree in computer science in 2005, all from the Massachusetts Institute of Technology (MIT), Cambridge, MA.

He is an Assistant Professor in the College of Computing at Georgia Tech, Atlanta. His research focuses on many aspects of computer networking and networked systems, including the design, measurement, and analysis of network protocols

and architectures, network management, network security, and anonymous communication systems.



Ramesh Johari (M'05) received the A.B. degree in mathematics from Harvard University, Cambridge, MA, in 1998, the Certificate of Advanced Study in Mathematics from the University of Cambridge, Cambridge, U.K., in 1999, and the Ph.D. degree in electrical engineering and computer science from the Massachusetts Institute of Technology (MIT), Cambridge, MA, in 2004.

He is currently an Assistant Professor of management science and engineering, and by courtesy, electrical engineering and computer science, at Stanford

University, Stanford, CA.



Hari Balakrishnan received the B.Tech. degree from the Indian Institute of Technology, Madras, in 1993, and the Ph.D. degree in computer science from the University of California at Berkeley in 1998.

He is a Professor in the Electrical Engineering and Computer Science Department and a member of the Computer Science and Artificial Intelligence Laboratory (CSAIL) at Massachusetts Institute of Technology (MIT), Cambridge, MA. His research interests are in the area of networked computer systems,

spanning overlay and peer-to-peer networks, Internet routing, congestion control, and architecture, wireless and sensor networks; network security, and distributed data management.